

Active Neural SLAM

Tuesday, 20 October 2020 11:34

Active Neural SLAM: (ANS)

- a modular and hierarchical approach to learn policies for exploring 3D environments.
- Uses analytical path planners with learned θ SLAM module, global & local policies.
- In the SLAM module: learning provides flexibility w.r.t. input modalities.
- In global policies: learning leverages structural regularities of the world.
- In local policies: learning provides robustness to errors in state estimation.
- Navigation problems:
 1. point goal tasks: navigating to specific coordinates.
 2. semantic navigation: finding the path to a specific scene or object.
- A core problem for navigation in unknown environments is exploration.

Learning exploration policies for navigation [Chen'19];

used end-to-end learning

pros: 1. flexibility w.r.t. input modalities (use RGB images)

2. robustness to errors in state estimation.

3. leverage structural regularities of the real world

→ more efficient behavior in previously unseen environment.

cons: expensive path-planning (uses imitation learning).

Solution:

learning (in the 3 pros) happens at different time scales
→ can be factored out.

- Learning-based navigation techniques: NN policies that reason via spatial representations

Neural Map: structured memory for DRL [Parisotto'18]

Neural SLAM [Zhang'17]

① topological representations

② Episodic curiosity through reachability [Savinov'19]

③ differential and trainable planners

Gated path planning networks [Lee'18]

Task setup:

the objective: maximizing the coverage in a fixed time budget.

the coverage: the total area in the map known to be traversable.

Habitat simulator [Savva'19]

- Actuation and sensor noise model:

agent pose (x, y, θ) , control command $\Delta u_a = (x_a, y_a, \theta_a)$

- The actuation noise E_{act} : is the difference between the actual agent pose P_t after the action and the intended agent pose $(P_0 + \Delta u)$

$$E_{act} = P_t - (P_0 + \Delta u) = (x^* - x_a, y^* - y_a, \theta^* - \theta_a)$$

- The sensor noise E_{sen} : is the difference between the sensor pose estimation P_t' and the actual pose P_t

$$E_{sen} = P_t' - P_t = (x^* - x^*, y^* - y^*, \theta^* - \theta^*)$$

- Actions (control commands):

Forward: $u_F = (0.25, 0, 0)$

Turn right: $u_R = (0, 0, -10^\circ \pi/180)$

Turn left: $u_L = (0, 0, 10^\circ \pi/180)$

- For each action, a separate Gaussian Mixture Model **GMM** is fitted for the actuation and sensor noise, making the total 6 models.

Methods:

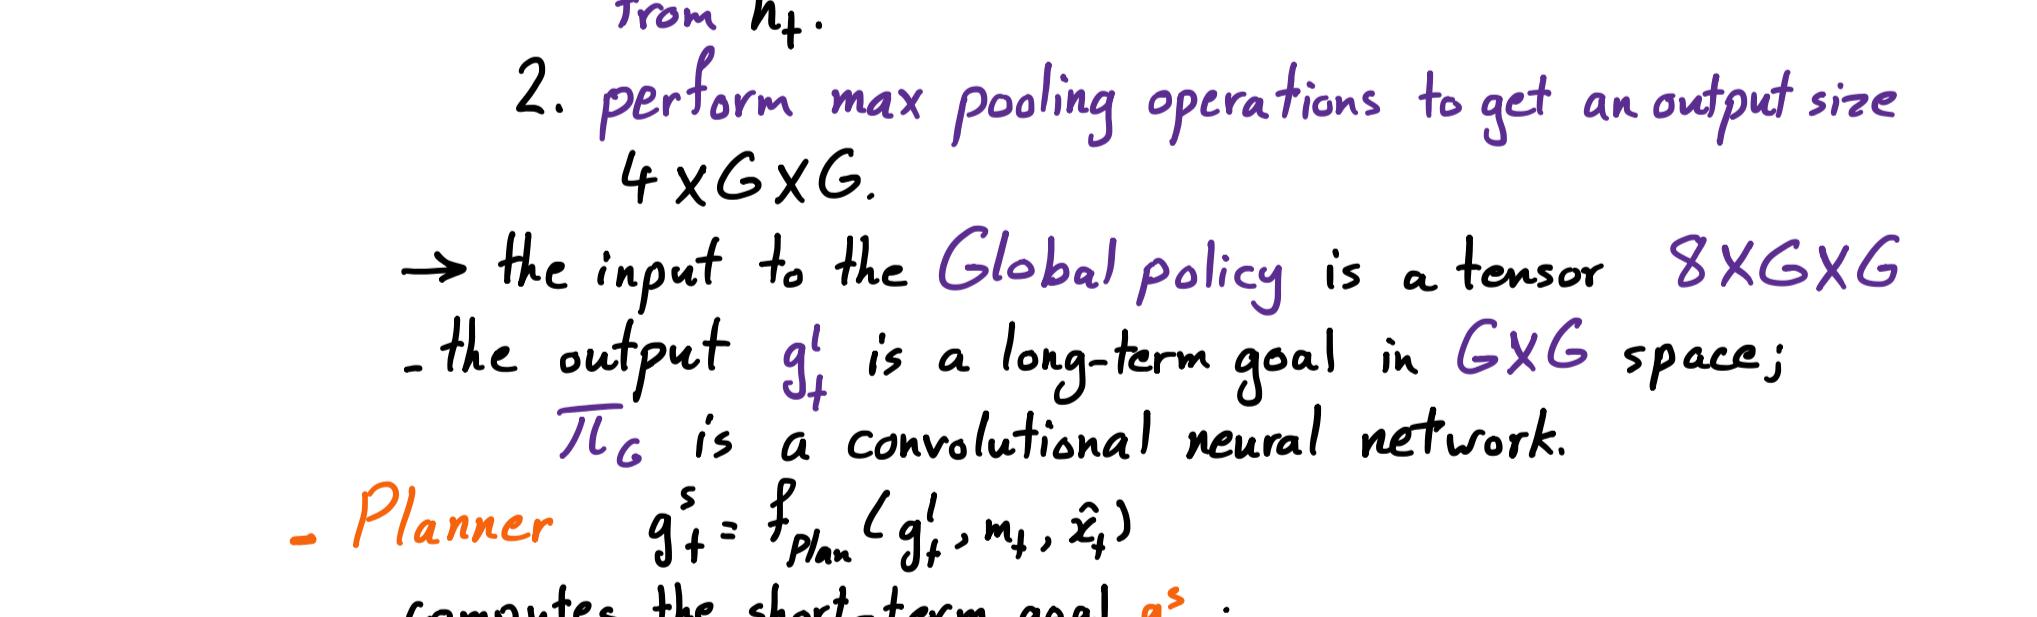


Figure 1: Overview of our approach. The Neural SLAM module predicts a map and agent pose estimate from incoming RGB observations and sensor readings. This map and pose are used by a Global policy to output long-term goal, which is converted to a short-term goal using an analytic path planner. A Local Policy is trained to navigate to this short-term goal.

- the spatial map has two channels:

1. denotes the probability of an obstacle at the corresponding location.
2. denotes the probability of that location being explored.

(explored = known to be free space or an obstacle.)

- Neural SLAM module:

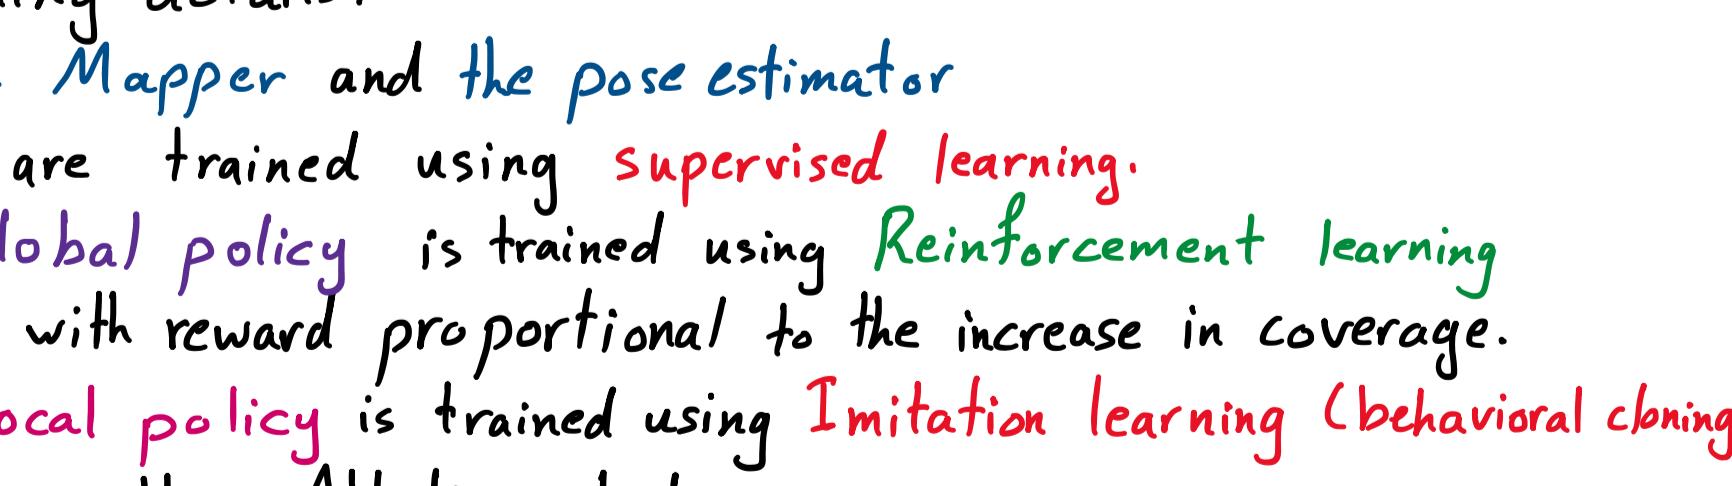


Figure 2: Architecture of the Neural SLAM module: The Neural SLAM module (f_{Map}) takes in the current RGB observation, s_t , the current and last sensor reading of the agent pose $x_{t-1:t}$, last agent pose estimate, \hat{x}_{t-1} and the map at the previous time step m_{t-1} and outputs an updated map, m_t and the current agent pose estimate, \hat{x}_t . 'ST' denotes spatial transformation.

$$m_t, \hat{x}_t = f_{SLAM}(s_t, x_{t-1:t}, \hat{x}_{t-1}, m_{t-1} | \theta_s)$$

- the trainable parameters θ_s consists of:

1. the mapper f_{Map} :

outputs a egocentric top-down 2D spatial map

$$p_t^{ego} \in [0, 1]^{2 \times V \times V} \quad (V \text{ is the vision range})$$

predicts the obstacles and the explored area in the current observation.

2. the pose estimator f_{PE} :

predicts the agent pose \hat{x}_t based on ① past pose estimate \hat{x}_{t-1} and ② the last two egocentric map predictions p_{t-1}^{ego}

- predicts the pose change by comparing maps after transforming the last to the current frame.

- Global policy $g_t^l = \pi_G(h_t | \theta_G)$

Its input $h_t \in [0, 1]^{4 \times M \times M}$ has 4 channels:

1-2. the spatial map m_t given by SLAM.

3. current position estimation \hat{x}_t given by SLAM.

4. visited locations $\hat{x}_t^k : k \in \{0, \dots, t\}$

two transformations performed on h_t :

1. subsample a window of size $4 \times G \times G$ around the agent from h_t .

2. perform max pooling operations to get an output size $4 \times G \times G$.

→ the input to the Global policy is a tensor $8 \times G \times G$

- the output g_t^l is a long-term goal in $G \times G$ space;

π_G is a convolutional neural network.

- Planner $g_t^s = f_{Plan}(g_t^l, m_t, \hat{x}_t)$

- g_t^s is transformed into relative distance and angle from the agent's

location + s_t RGB observations → π_L

- π_L is an RNN with a pretrained ResNet18 as visual encoder.

- Training details:
- The Mapper and the pose estimator are trained using supervised learning.
- Global policy is trained using Reinforcement learning with reward proportional to the increase in coverage.
- Local policy is trained using Imitation learning (behavioral cloning).
- From results: Ablation study
- Local policy overcomes false positives encountered in mapping.
- Global policy ignores small spaces and choose distant long-term goals → higher coverage.
- ANS reaches performance $0.789 / 0.703$ SPL / Succ in 1 million frames
- Succ: success rate.
- SPL: success weighted by path length (normalized inverse).
- ANS $> 75 \times$ speed-up as compared to best RL baseline.

Architecture:

f_{map} : ResNet 18 convolutional layers → embedding

→ 2 fully connected layers

→ 3 deconvolutional layers → first-person top-down 2D map

binary cross-entropy loss $\times 1$

f_{PE} : 3 convolutional layers

→ 3 fully connected layers

MSE loss $\times 1000$

π_G : PPO with reward = increase of coverage $m^2 \times 0.02$

π_L : binary cross-entropy loss.